



SDU

Deep Learning

Fall
2021

Introduction to KERAS 2

Three API Styles

- **The Sequential Model**
 - Dead simple
 - Only for single-input, single-output, sequential layer stacks
 - Good for 70+% of use cases
- **The functional API**
 - Like playing with Lego bricks
 - Multi-input, multi-output, arbitrary static graph topologies
 - Good for 95% of use cases
- **Model subclassing**
 - Maximum flexibility
 - Larger potential error surface

The Functional API

```
from tensorflow.keras.models import Model
from tensorflow.keras.layers import Dense, Input

inputs = Input(shape=(10,))
x = Dense(64, activation='relu')(inputs)
x = Dense(64, activation='relu')(x)
output = Dense(10, activation='softmax')(x)

Model = Model(inputs, output)
Model.compile(optimizer='rmsprop',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
model.fit(x, y, epochs=10, batch_size=32)
```

Defining the Model

```
from tensorflow.keras import models
from tensorflow.keras import layers

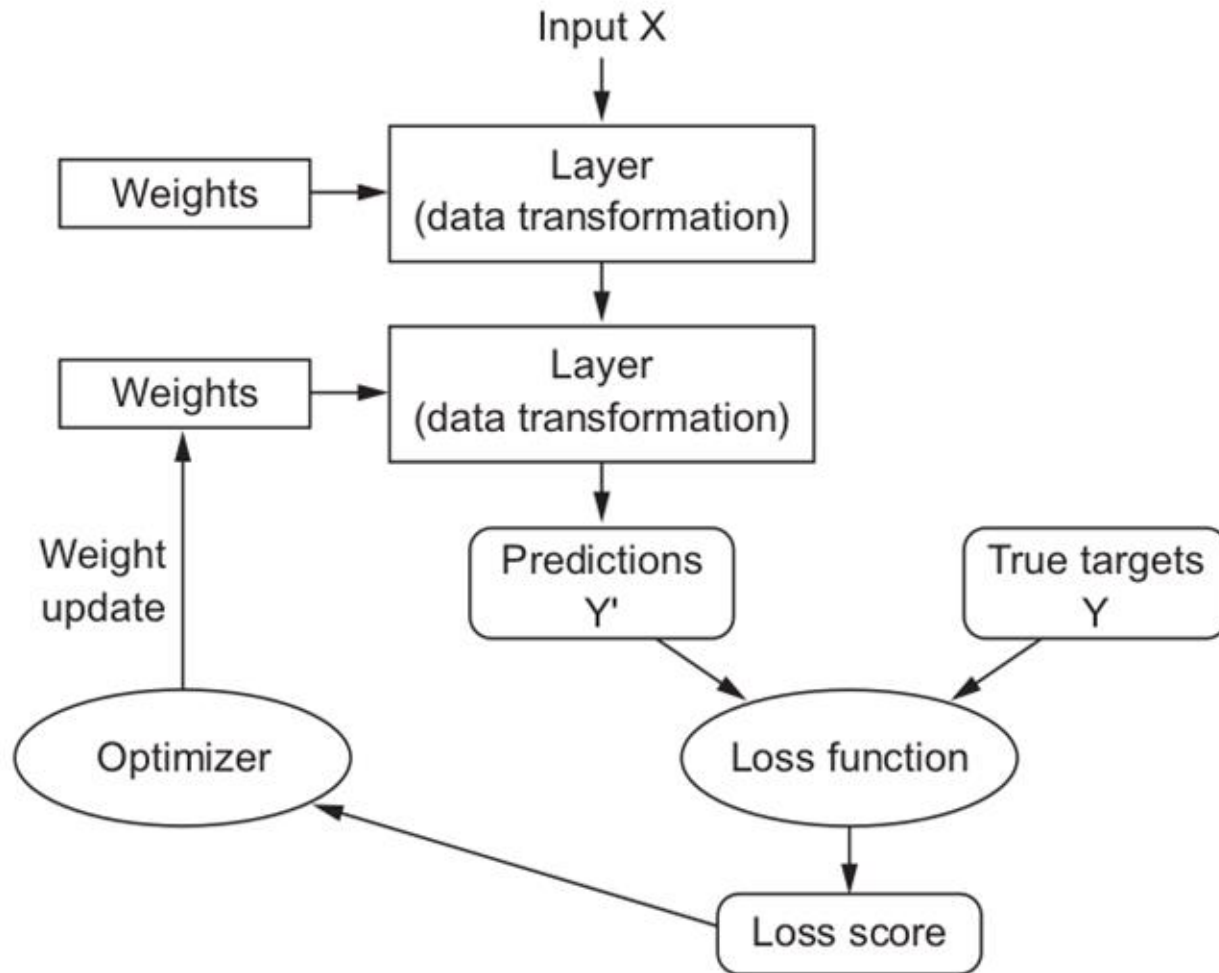
model = models.Sequential()
model.add(layers.Dense(32, activation='relu', input_shape=(784,)))
model.add(layers.Dense(10, activation='softmax'))
```

- The same model defined using the functional API:

```
input_tensor = layers.Input(shape=(784,))
x = layers.Dense(32, activation='relu')(input_tensor)
output_tensor = layers.Dense(10, activation='softmax')(x)

model = models.Model(inputs=input_tensor, outputs=output_tensor)
```

Alrighty, let's put it together



Regularization

- Adding weight regularization

```
from tensorflow.keras import regularizers
```

```
model = models.Sequential()
```

```
model.add(layers.Dense(16, kernel_regularizer=regularizers.l2(0.001),  
                        activation='relu', input_shape=(10000,)))
```

```
model.add(layers.Dense(1, activation='sigmoid'))
```

- `l2(0.001)` means every coefficient in the weight matrix of the layer will add $0.001 * \text{weight_coefficient_value}$ to the total loss of the network.
- Note that because this penalty is only added at training time, the loss for this network will be much higher at training than at test time.

Dropout

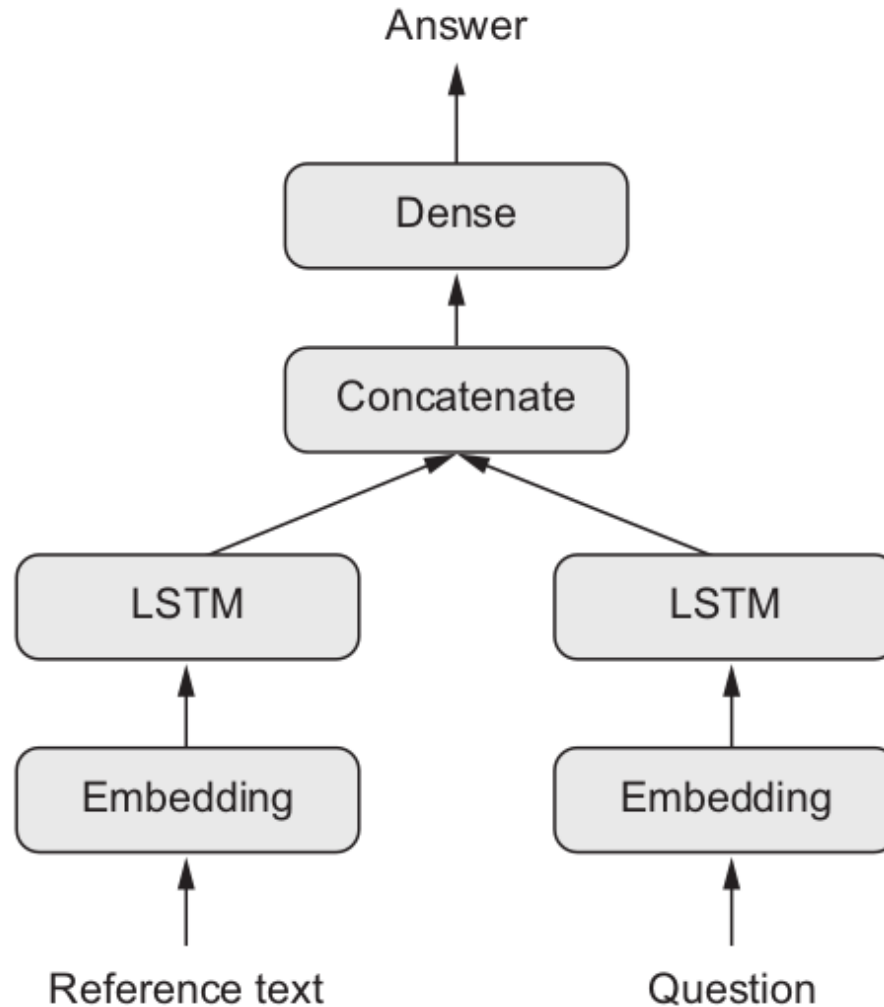
```
model = models.Sequential()

model.add(layers.Dense(16, activation='relu', input_shape=(10000,)))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(16, activation='relu'))
model.add(layers.Dropout(0.5))

model.add(layers.Dense(1, activation='sigmoid'))
```

Multi Input Models



Multi Input Model using the functional API

```
text_input = Input(shape=(None,), dtype='int32', name='text')
embedded_text = layers.Embedding(64, text_vocabulary_size)(text_input)
encoded_text = layers.LSTM(32)(embedded_text)

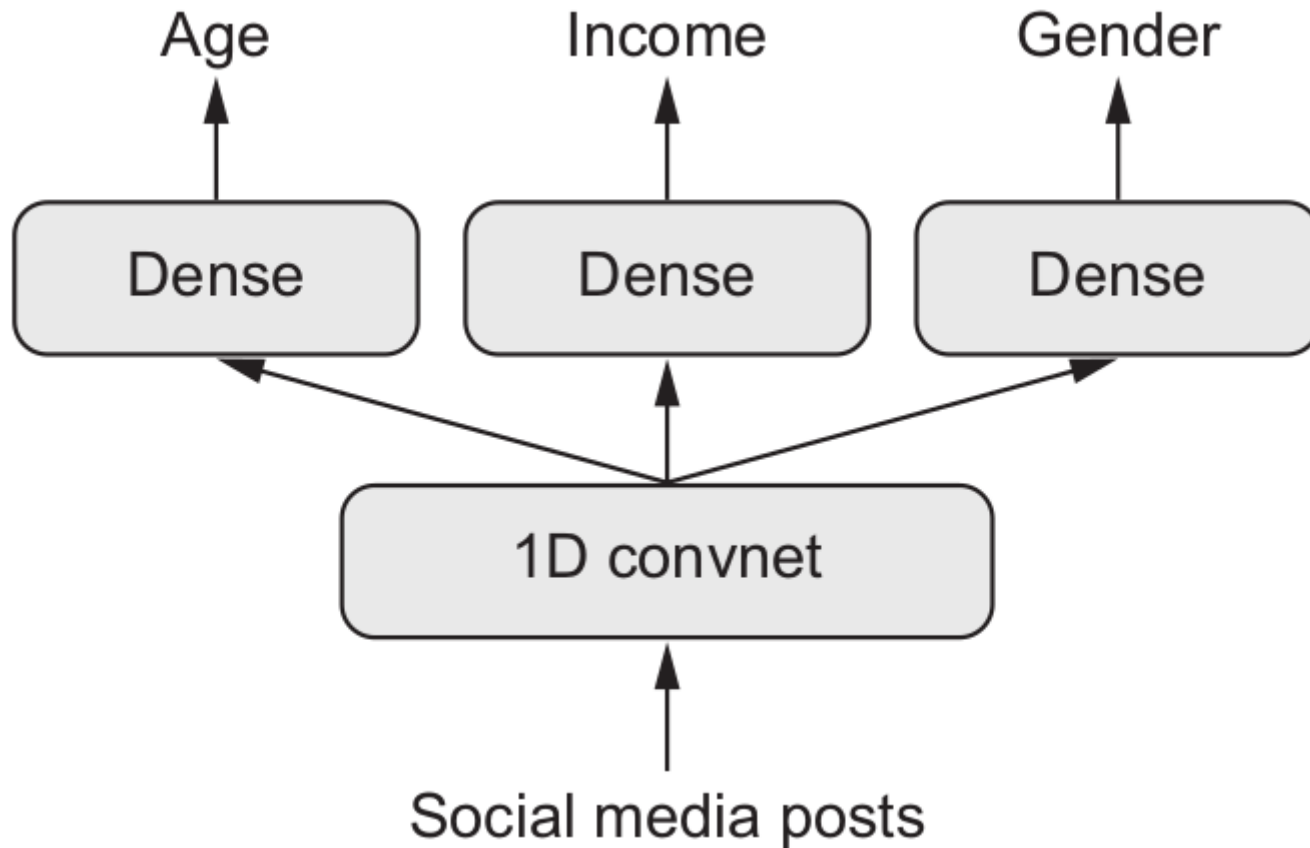
question_input = Input(shape=(None,), dtype='int32', name='question')
embedded_question = layers.Embedding(32, question_vocabulary_size)(question_input)
encoded_question = layers.LSTM(16)(embedded_question)

concatenated = layers.concatenate([encoded_text, encoded_question], axis=-1)
answer = layers.Dense(answer_vocabulary_size, activation='softmax')(concatenated)

model = Model([text_input, question_input], answer)

model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['acc'])
```

Multi-Output Models



|

Multi-Output Models

```
posts_input = Input(shape=(None,), dtype='int32', name='posts')
embedded_posts = layers.Embedding(256, vocabulary_size)(posts_input)
x = layers.Conv1D(128, 5, activation='relu')(embedded_posts)
x = layers.MaxPooling1D(5)(x)
... # Construct the network how you like it
x = layers.Dense(128, activation='relu')(x)

age_prediction = layers.Dense(1, name='age')(x)
income_prediction = layers.Dense(num_income_groups, activation='softmax',
                                name='income')(x)
gender_prediction = layers.Dense(1, activation='sigmoid', name='gender')(x)

model = Model(posts_input,[age_prediction, income_prediction, gender_prediction])
```

Multi-Output Models

```
model.compile(optimizer='rmsprop',  
              loss=['mse', 'categorical_crossentropy', 'binary_crossentropy'])
```

```
model.compile(optimizer='rmsprop',  
              loss={'age': 'mse',  
                    'income': 'categorical_crossentropy',  
                    'gender': 'binary_crossentropy'})
```

- Fitting, etc, of the model remains the same as with a normal network.

Callbacks

- When training a model, there are many things you cannot predict
- Sometimes it would be helpful to intervene when something goes wrong
- Keras provides callbacks:
 - Model checkpointing: Saving the current weights of the model at different points during training.
 - Early stopping: Interrupting training when the validation loss is no longer improving (and of course, saving the best model obtained during training).
 - Dynamically adjusting the value of certain parameters during training: Such as the learning rate of the optimizer.
 - Logging training and validation metrics during training, or visualizing the representations learned by the model as they're updated: The Keras progress bar is a callback!

Implementing your own Callback

- Callbacks are implemented by sub-classing the class `keras.callbacks.Callback`.
- You can implement the following functions:
 - `on_epoch_begin`
 - `on_epoch_end`
 - `on_batch_begin`
 - `on_batch_end`
 - `on_train_begin`
 - `on_train_end`

TensorBoard / Weights and Biases

- Callbacks allow you to send run information to services that help you track and visualize runs.
- TensorBoard runs locally on a logs folder
 - Highly modular, but base features are lacking
- Weights and Biases runs in the cloud by sending the information to their servers. Only need to log onto their servers once on your local machine
 - No/Low modularity, but base features are rich and easy to setup

TensorBoard

```
tensorboard_callback = TensorBoard(log_dir='/path/to/logs/', update_freq='batch')  
model.fit(x=train_data, y=val_data, epochs=10, callbacks=[tensorboard_callback])
```

Terminal or CMD:

```
C:\Users\Name\Desktop>tensorboard --logdir /path/to/logs/  
>> 'TensorBoard 2.4.0 at http://localhost:6006/ (Press CTRL+C to quit)'
```


Weights and Biases

```
import wandb
wandb.init(project='project name')

from wandb.keras import
model.fit(x=train_data, y=val_data, epochs=10, callbacks=[WandbCallback()])
```

Early Stopping

```
import tensorflow.keras as keras

callbacks_list = [
    keras.callbacks.EarlyStopping(
        monitor='acc',
        patience=1),
    keras.callbacks.ModelCheckpoint(
        filepath='my_model.h5',
        monitor='val_loss',
        save_best_only=True)]

model.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])

model.fit(x, y, epochs=10, batch_size=32, callbacks=callbacks_list,
        validation_data=(x_val, y_val))
```

Predictions

```
model = load_model('model.h5')
predictions = model.predict(test_data)

# Sigmoid (Boolean)
predictions = (predictions > 0.5).astype(int).reshape(test_labels.shape)

# Softmax (Multi class)
predictions = predictions.argmax(axis=-1)

wrong_data = test_data[test_labels != predictions]
wrong_preds = predictions[test_labels != predictions]
```

Data Generators

- Data generators allows you to import, train, and discard batches of data from memory.
- Provides the option for 'on the fly' data augmentation
- Useful for large datasets, which cannot be stored in memory

Data Generators

```
tf.keras.preprocessing.image.ImageDataGenerator(  
    featurewise_center=False,  
    samplewise_center=False,  
    featurewise_std_normalization=False,  
    samplewise_std_normalization=False,  
    zca_whitening=False,  
    zca_epsilon=1e-06,  
    rotation_range=0,  
    width_shift_range=0.0,  
    height_shift_range=0.0,  
    brightness_range=None,  
    shear_range=0.0,  
    zoom_range=0.0,  
    channel_shift_range=0.0,  
    fill_mode="nearest",  
    cval=0.0,  
    horizontal_flip=False,  
    vertical_flip=False,  
    rescale=None,  
    preprocessing_function=None,  
    data_format=None,  
    validation_split=0.0,  
    dtype=None)
```

➤ <https://keras.io/api/preprocessing/image/>

Data Generators

```
train_datagen = ImageDataGenerator(  
    rescale=1./255, shear_range=0.2, zoom_range=0.2, horizontal_flip=True)  
test_datagen = ImageDataGenerator(rescale=1./255)  
  
train_generator = train_datagen.flow_from_directory(  
    'data/train', target_size=(150, 150), batch_size=32, class_mode='binary')  
validation_generator = test_datagen.flow_from_directory(  
    'data/validation', target_size=(150, 150), batch_size=32, class_mode='binary')  
  
model.fit(train_generator, steps_per_epoch=2000, epochs=50, v  
    alidation_data=validation_generator, validation_steps=800)
```

➤ <https://keras.io/api/preprocessing/image/>

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset

- What will your input data be?
- What type of problem are you facing? Classification? Regression?
- ...

1. Choosing a measure of success
2. Decide on the evaluation protocol
3. Preparing your data
4. Define a model better than base-line
5. Scaling up: Make the model overfit
6. Regularizing your model and tuning your hyperparameters
7. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
- 2. Choosing a measure of success**
 - How do you measure if the model is successful
 - Not to be confused with the loss function which is often only a surrogate for what you actually want achieve
1. Decide on the evaluation protocol
2. Preparing your data
3. Define a model better than base-line
4. Scaling up: Make the model overfit
5. Regularizing your model and tuning your hyperparameters
6. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
 2. Choosing a measure of success
 3. **Decide on the evaluation protocol**
 - hold-out validation
 - Cross-validation
-
1. Preparing your data
 2. Define a model better than base-line
 3. Scaling up: Make the model overfit
 4. Regularizing your model and tuning your hyperparameters
 5. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. **Preparing your data**
 - Clean data
 - Normalize data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
 2. Choosing a measure of success
 3. Decide on the evaluation protocol
 4. Preparing your data
 5. **Define a model better than base-line**
 - Last-layer activation
 - Loss function
 - Optimization Algorithm
-
1. Scaling up: Make the model overfit
 2. Regularizing your model and tuning your hyperparameters
 3. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
 2. Choosing a measure of success
 3. Decide on the evaluation protocol
 4. Preparing your data
 5. Define a model better than base-line
 6. **Scaling up: Make the model overfit**
 - Add layers.
 - Make the layers bigger.
 - Train for more epochs.
-
1. Regularizing your model and tuning your hyperparameters
 2. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. **Regularizing your model and tuning your hyperparameters**
 - This will take the most time
 - Add dropout.
 - Try different architectures: add or remove layers.
 - Add L1 and/or L2 regularization
8. Finalize your final model

Deep Learning Workflow

1. Defining the problem and assembling a dataset
2. Choosing a measure of success
3. Decide on the evaluation protocol
4. Preparing your data
5. Define a model better than base-line
6. Scaling up: Make the model overfit
7. Regularizing your model and tuning your hyperparameters
8. **Finalize your final model**
 - Save and distribute the model